

# SOHAN ANAND HANAGANDI

AI Product Manager · Builder

+1(857) 370-2511 | [sohananandhanagandi@gmail.com](mailto:sohananandhanagandi@gmail.com) | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

## SUMMARY

AI PM, 5 yrs across B2B SaaS, gaming, AI. Shipped Multi-Agent Orchestration + RAG at Fabrix.ai for Tata Communications + 2 Fortune 500 accounts. Scaled PlayStation Dev Portal 5x. Now building multi-agent systems & Agentic AI Agents solo, end to end.

## EDUCATION

Northeastern University, MS, Project Management (Product & Technology Focus)  
Gogte Institute of Technology, BE, Computer Science

Sep 2023 – Dec 2025  
Aug 2016 – Mar 2020

## EXPERIENCE

**Product Manager** | Community Dreams Foundation | San Francisco Bay Area, CA Feb 2026 – Present

- Shipped intake, triage, & document-generation AI agents to production across 2 Legacy products, replacing manual ops workflows with LLM pipelines. Owned the full stack from prompt engineering & pipeline orchestration to evaluation and monitoring
- Owning eval harness per agent (golden-set, LLM-as-judge); running blind comparisons against staff baselines before each rollout.
- Aligned 3 stakeholder groups on quarterly AI roadmap; unlocked engineering capacity for 2 priority initiatives.

**AI Product Co-op** | Fabrix.ai | Pleasanton, CA Jan 2025 – Jul 2025

- Drove multi-agent orchestration for Tata Communications + 2 Fortune 500 accounts; cut onboarding ~10%.
- Shipped RAG pipelines (custom chunking, reusable retrieval); lifted golden-set eval relevance ~10pp.
- Replaced GPT-4 with fine-tuned open-source LLM in prod; held eval quality flat while cutting per-token cost.
- Caught at-risk enterprise accounts pre-value via telemetry; reprioritized backlog and shipped 4 fixes that closed renewal.

**Product Manager** | Sony Interactive Entertainment, PlayStation Dec 2022 – Sep 2023

- Led 14-person cross-functional team scaling PlayStation Developer Portal from 800 to 4,000+ partners in 10 months.
- Diagnosed funnel drop-off in dev-portal onboarding; redesigned compliance workflow, cutting partner activation ~22%.
- Built SQL-backed ROI model that secured VP approval to expand into 2 regional markets.

**Product Manager** | Qlik Technology Jan 2020 – Nov 2022

- Killed 2 features (\$250K pipeline) after <5% cohort adoption; redirected 4 engineers to retention, recovering \$285K ARR.
- Caught a 90-day churn spike in SQL/Mixpanel; shipped 2 targeted interventions, recovering \$180K ARR over 3 quarters.
- Ran vendor eval across 8 tools; surfaced 3 gaps that reshaped the roadmap and unlocked 2 enterprise wins (\$95K ACV).

## PROJECTS

**Trip Booking Concierge – Autonomous AI Agent with Trust-by-Default Approval Gate** | [Live Demo](#) May 2026

- Defined 0-to-1 product vision for autonomous AI trip-booking agent eliminating \$2K+ irreversible booking errors for travelers.
- Spec'd two-step approval gate balancing UX friction vs. financial safety, cutting unsafe-booking risk to 0 across adversarial tests.
- Delivered 100% fact provenance and zero approval bypasses across 12+ red-team scenarios, validating safe-by-default agent UX.

**Kira AI: Production RAG AI Voice Agent** | [Live on Portfolio](#) Feb 2026

- Built an AI agent so recruiters get answers about my work in 2 minutes; replaces the screening calls I used to take.
- Engineered pgvector schema on Supabase + OpenAI embeddings + Claude Haiku streaming with context windowing & refusal logic.
- Deployed with prompt guardrails and refusal logic for out-of-scope queries; cut screening overhead to zero, live on Vercel.

**BrandBrief AI: RAG Template Matching Engine** | [Live App](#) | [Case Study](#) Jan 2026

- Diagnosed brief-to-template matching as conversion blocker; shipped Claude + pgvector RAG engine, processing 4,500+ records.
- Replaced manual ops overhead with no-LangChain RAG on pgvector + Supabase; delivered <3s P50 time with Claude inference.
- Achieved ~70% top-1 retrieval accuracy, removing hours of manual search; moved 4,500+ templates to live production with auth.

## SKILLS

**AI / ML** : LLMs, Agentic Workflows, Multi-Agent Orchestration, RAG, Fine-Tuning (SFT), Vector Embeddings, Function Calling, Structured Outputs, LangChain, LangGraph, Prefix Caching, Guardrails, Voice Agents, MCP, Evals

**AI Dev Stack** : Claude Code, Cursor, Anthropic / OpenAI APIs, pgvector, Supabase, LLM-as-Judge & Golden-Set Evals, Prompt Engineering, Streaming Inference

**Product & Strategy** : 0→1 Product, PRDs, Backlog Prioritization, Roadmapping, A/B Testing, Release Management, OKRs, Cohort & Funnel Analysis, Activation & Retention, SaaS Metrics (ARR, NPS, Churn), Agile/Scrum

**Data & Tooling** : SQL, Python (FastAPI, requests), TypeScript/React, Mixpanel, Amplitude, Productboard, Figma, Jira, GitHub